

SHASHANK YADAV

SENIOR JAVA BACKEND ENGINEER | DISTRIBUTED SYSTEMS | DATA ENGINEERING | KAFKA | SDE-3

New Delhi, India | +91-8218233193 | shashankyadav.1299@gmail.com | [LinkedIn](#) | [GitHub](#) | [Articles](#)

PROFESSIONAL SUMMARY

Senior Java Backend Engineer with 5+ years of experience designing and developing Java/Spring Boot microservices, REST APIs, Kafka-based event-driven systems, and distributed applications. Strong expertise in Java 8/11/21, Spring Boot, Spring Security, JPA/Hibernate, Apache Kafka, Kafka Streams, SQL/Oracle, and cloud-based applications. Experienced in backend architecture, asynchronous processing, caching, performance optimization, production troubleshooting, and automated testing. Proven track record of leading end-to-end technical initiatives, resolving P0/P1 production issues, and driving solutions from design through production delivery.

TECHNICAL SKILLS

Languages: Java 8/11/21, Python, SQL, PL/SQL, JavaScript

Java & Backend: Java Streams, Multithreading, Concurrency, CompletableFuture, OOP, Design Patterns, Spring Boot, Spring MVC, Spring Security, Spring Data JPA, Hibernate, REST APIs, Microservices, Apache Camel, JMS

Distributed Systems: Apache Kafka, KStream, Consumer Groups, Partitions, Offset Management, Event-Driven Architecture, Retry/DLQ, Avro, Schema Registry

Architecture: Microservices Architecture, API Gateway, Distributed Transactions, Circuit Breaker, System Design

Databases: Oracle, SQL Server, Snowflake, SQL/PLSQL, Query Optimization

Cloud & DevOps: AWS, Kubernetes, Jenkins, GitLab CI/CD, Maven/Gradle

Testing & Observability: JUnit, Mockito, K6, Locust, Kibana, OpenObserve

PROFESSIONAL EXPERIENCE

HashedIn by Deloitte - Software Engineer 2 | May 2024 - Present

Data Engineering | April 2026 - Present

- Designed and implemented 14 real-time data pipelines using Java, Apache Kafka and Kafka Streams/KStream, processing approximately 50,000 messages/hour across 5 integrated systems.
- Built Java-based Kafka producer and consumer applications for event-driven data movement and processing across interconnected systems.
- Implemented end-to-end XML-to-Avro transformation, including serialization and deserialization for standardized downstream processing.
- Integrated Confluent Schema Registry and standardized Avro message structures across pipelines.
- Processed transformed data into Snowflake tables for downstream analytics and reporting.
- Designed the end-to-end architecture for real-time data pipelines, decomposing the solution into independently deployable components and driving technical design, stakeholder alignment, and production delivery.
- Designed reliable cross-system data transformation and movement workflows, including handling of missing-entity scenarios.

Backend Engineering | May 2024 - March 2026

- Designed and implemented a Kafka-based comparator module in Java from scratch to compare legacy and upgraded microservice responses, identifying missing, additional, and incorrect fields during the gateway migration.
- Took end-to-end ownership of gateway migration for multiple tenants, delivering 30-35 enhancements, fixes and configuration changes through development, testing, deployment and production stabilization.
- Implemented Kafka RO event payload changes across 3 services to standardize asynchronous inter-module communication within the international Returns platform.
- Enhanced an automated refund reprocessing pipeline with 5-6 improvements and 8 new error-handling scripts, supporting approximately 1,000 orders/day and saving 2-3 engineering hours/day.
- Developed automation and database utilities to insert and correct 15,000+ production seller-address records, resolving a critical data-integrity issue affecting return-order creation.
- Developed automated jobs, scripts, alerts and failure-reprocessing workflows that reduced manual operational effort by approximately 95-98% across targeted processes.
- Resolved complex transactional and order-placement APIs, including P0/P1 production incidents, while investigating logs and downstream system behavior.
- Owned the IQS application-to-application authentication feature from design through production deployment as a critical dependency for downstream services.
- Established production monitoring and endpoint failure alerting using OpenObserve, improving visibility into P0/P1 issues.
- Led JUnit testing and code reviews, strengthening validation and reliability across existing production use cases.
- Partnered with US client teams and product stakeholders on technical requirements, solution design, demonstrations, and delivery refinement.

Accenture | Aug 2021 – Apr 2024

Custom Software Engineering Analyst | Sep 2022 - Apr 2024

- Improved legacy microservice performance by 35% through caching, improving application responsiveness under production load.
- Designed, developed and resolved production issues in transactional APIs for payments and order placement, including recursive seat-reassignment logic.

- Executed Oracle SQL/PLSQL scripts to resolve data synchronization discrepancies across schemas, correcting issues affecting approximately 100,000 records.
- Led design, development, testing and validation of production backend services while resolving critical production issues end-to-end.
- Developed and executed 150+ integration tests across services to improve production readiness and reliability.
- Implemented Kafka RO event payload properties across 3 repositories, improving consistency of event data across modules.

Application Development Associate / Aug 2021 - Aug 2022

- Upgraded 30+ services within two months to implement client-requested functionality across existing applications.
- Designed and implemented 100+ unit tests, improving code quality, reliability and maintainability.
- Prioritized and resolved 50+ critical and non-critical production defects while meeting quality and delivery requirements.
- Prepared technical documentation covering requirements, designs and feature implementations.

OPEN SOURCE PROJECTS

Kafka Viewer - Python, Kafka, Streamlit, Confluent Schema Registry

- Built and published a lightweight Python package on PyPI for local Kafka inspection, supporting Kafka security configurations, SASL/SSL, Java keystore/truststore conversion to PEM, Schema Registry, and Confluent Avro deserialization.
- Implemented topic statistics, message filtering, consumer-group-based inspection, and read-only Kafka consumption without committing or modifying consumer offsets.

EDUCATION

Bachelor of Technology - Computer Science and Engineering | Galgotias University, Greater Noida | 2017 - 2021 | CGPA: 8.58/10

CERTIFICATIONS

- AWS Certified AI Practitioner
- Apollo Graph Developer Associate Certification

PROFESSIONAL DEVELOPMENT

- Java Upskilling - Internal Training, 2026
- Data Engineering Java Capability Training - 2026
- Databricks - Hands-on Training / Project Experience

AWARDS & LEADERSHIP

- Excellence Award - HashedIn by Deloitte
- Accenture Celebrates Excellence Award
- Led a project from design through production delivery, coordinating a team of 4 engineers and driving execution to successful completion.
- Conducted 40+ Java/Python technical interviews across lateral hiring and campus drives.
- Mentored engineers preparing for technical interviewer roles.
- Served as GenAI training Track Helper and Evaluator.
- Delivered a session on automated refund processing using Looper Pro.